

UNIVERSITÀ
DEGLI STUDI
DI PADOVA



DIPARTIMENTO
DI INGEGNERIA
DELL'INFORMAZIONE

DEPARTMENT OF INFORMATION ENGINEERING
BACHELOR'S DEGREE IN COMPUTER ENGINEERING

Fusing vision and inertial measurements for autonomous
navigation in narrow channels

Supervisor:
Prof. Damiano Varagnolo

Candidate:
Thomas Sanavia

ACADEMIC YEAR: 2024/2025
Graduation date: — — —-

A ...

Quote
Author

Contents

Introduction	1
1 Introduction	3
1.1 Autonomous Surface Vehicles: challenges and applications . .	3
1.2 Attitude Estimation: roll, pitch, yaw	5
1.3 Computer Vision for Localization and Attitude Estimation . .	7
1.4 Fiducial markers and ArUco tags	9
1.4.1 Sub-section title	11
Bibliografia	13

Abstract

This thesis addresses the design, development, and validation of an automatic system for detecting roll, pitch, and yaw of an autonomous boat. The main objective is to create a low-cost, low-power, and easily transportable solution capable of accurately determining the boat's attitude. To achieve this, both computer vision and an onboard inertial sensor will be used. By fusing data from these two sources, a precise estimate of the boat's attitude will be made available to the entire system. For the computer vision part, two cameras will be used to exploit stereo vision. The attitude estimation will be processed on a Raspberry Pi 5, chosen for its low cost. For this purpose, ArUco Tags—two-dimensional fiducial markers commonly used in robotics and augmented reality for pose estimation—will be employed. Regarding the inertial sensor, the onboard device will provide velocities and accelerations along the three axes. Finally, an extended Kalman filter, which will utilize both inertial and computer vision data, will be applied to reduce drift and improve the accuracy of roll, pitch, and yaw estimation.

Introduction

In recent years, the field of **autonomous systems** has played an increasingly central role in robotics. All this has also been made possible thanks to solutions such as **Robot Operating System 2 (ROS2)**, which is one of the platforms for the development of robotic applications. This has been achieved thanks to its *modular architecture*, support for *real-time communication*, and its *open-source* nature, which has made it an excellent choice even for advanced robotic systems.

This thesis is part of the **Autodocking project**, focused on the ability of a boat to navigate and dock autonomously. Although the structure and the basic control systems were already in place, a **computer vision-based system** for autonomous navigation was missing. My contribution was related to the creation of a **ROS2 node**; more specifically, I dealt with the entire **pipeline** ranging from **image acquisition**, **attitude estimation**, and subsequent **sensor fusion** with the inertial sensor data. To achieve this, I needed a **ROS2 node** that handled the acquisition of images from the various cameras installed on the boat and their publication on a *topic*, so that my algorithm could access them.

Chapter 1

Introduction

1.1 Autonomous Surface Vehicles: challenges and applications

Unmanned Surface Vehicles (**USVs**) are particular surface vessels capable of navigating without the presence of a human crew to maneuver them. They can be *remotely operated* through radio or satellite communication, or be completely *autonomous* and controlled by a computer or artificial intelligence; in this case, they are referred to as **Autonomous Surface Vehicles (ASVs)**. ASVs aim to make journeys *safer, reduce costs, increase operational continuity*, and carry out *prolonged missions* in *complex* and even *risky* situations for humans.

All this has been made possible thanks to significant technological evolution in **sensors** (satellite, IMU, cameras, LIDAR, radar, etc.), major devel-

opments in **robotics libraries**, and onboard **computational power** that is increasingly greater and more affordable. Despite this, the **marine environment** presents major challenges. For example, the *plane of the sea surface* is not a fixed reference but changes independently of the boat’s movements; the surface creates *reflections* that can interfere with instrumentation; and *wind and waves* can put stability control systems to the test. In the marine context, therefore, **attitude estimation** (roll, pitch, and yaw) is essential, as it is decisive both for the *success of the journey* and for the *precision required in maneuvers*, such as docking.

The main challenges concern **sensors and algorithms**. The first challenge involves **navigation with GNSS** (Global Navigation Satellite System), which degrades in proximity to port infrastructures, bridges, or areas with many vertical obstacles, as part of the visible satellites are lost. This problem has been partly mitigated with **inertial sensors (IMU)**, which, however, suffer from another issue: *drift*. The second challenge is the **detection of possible obstacles** at sea, such as buoys, debris, or even other vessels, often under conditions of *reduced visibility* or *reflections*. The third challenge is **control and stability**, made complex by *wave motion* that often causes very rapid variations. The last challenge is the ability to make **real-time decisions**, including very quick changes dictated by sudden shifts in the operational situation.

Despite the criticality of these challenges, ASVs are used in various **applications**. In the **industrial field**, they are employed for *infrastructure*

inspection or *short-range logistical support*. In **defense and security**, they are used in *special operations* such as *explosive ordnance disposal*. In the **environmental field**, they are used for *monitoring environmental parameters* and *collecting meteorological data* over extended time horizons, thereby reducing risks for crews.

Within this framework, the present thesis focuses on one main issue: obtaining an **accurate real-time estimation** of the boat's attitude. To achieve this, **computer vision techniques** (stereo vision and ArUco markers), **inertial measurements** through the IMU, and an **Extended Kalman Filter** will be used to fuse the data, with the aim of *improving the reliability of the estimation*.

1.2 Attitude Estimation: roll, pitch, yaw

To describe the **attitude** of a vessel, the orientation of the rigid body with respect to a reference system is used. This is done through the three **Euler angles**[1]: *roll* (rotation on the longitudinal axis), *pitch* (rotation on the transverse axis), and *yaw* (rotation on the vertical axis). These three quantities are **crucial** as they influence **stabilization** and **trajectory control** and are fundamental in **precision maneuvers**.

Attitude estimation can be obtained from different sources, each with its own advantages and disadvantages. **Inertial sensors (IMU)**[2], through gyroscopes and accelerometers, provide roll and pitch with a **high update frequency**. However, these estimates are affected by *drift*, caused by the




Figure 1.1: Three rotational degrees of freedom of a ASV

accumulation of errors over time. For yaw, a **magnetometer** or an **external observation** is required, as the IMU does not provide a direct measurement of this angle. A **visual sensor** allows estimating the attitude through *known structures* (e.g., ArUco tags) or through the environment if it allows it (e.g., using the *horizon line*). This latter method is less affected by the drift problem but is highly sensitive to **environmental conditions** and also has a **lower update frequency**.

In the **marine context**, a robust pipeline is represented by the IMU, which provides a **high-frequency prediction** and, periodically, through **computer vision**, a measurement relative to the ArUco marker to reduce error accumulation. Through an **EKF filter**[3], the different contributions are integrated in order to obtain a **coherent estimate** even if the conditions are adverse or if one of the sources fails.

In summary, **attitude estimation** for an **ASV** requires a compromise between the **speed of the readings** and their **accuracy**. The **integration** between *inertial* and *visual* measurements provides a **reliable estimate** ca-

pable of supporting **autonomous navigation**. Subsequently, the different **implementation choices** to obtain the attitude estimation of the **Blue Boat** will be presented.

1.3 Computer Vision for Localization and Attitude Estimation

Computer vision is an essential source for estimating the attitude of the ASV. It provides **roll**, **pitch**, and **yaw** with respect to a known reference, namely the *ArUco marker*. In our case, it plays a *complementary role* with respect to the IMU, which provides an absolute measurement that is not obtained through temporal integration, but is instead sensitive to visual conditions (*reflections, backlighting, fog, etc.*).

The goal of this chapter is to outline the principles and choices that allow for reliable attitude estimation using a **stereo camera pair**, leveraging *ArUco markers* and a geometric calibration of the two cameras.

The pipeline is structured into four main steps:

- **Stereo camera calibration:** estimation of intrinsic parameters and distortions, followed by the calculation of rotation and translation between the two cameras.
- **Marker detection:** ArUco markers enable highly reliable identification, provided they are properly scaled according to the distance and have a sufficient pixel resolution.

- **Attitude estimation:** application of a pose estimation algorithm which, knowing the 2D corners and the 3D points of the pattern, can estimate the values of roll, pitch, and yaw.
- **Measurement validation:** a decision filter determines whether to accept the measurement, discard it, or reduce its influence.

Triangulation, made more reliable by **epipolar rectification**, allows for the estimation of depth around the marker. These depth cues are used as additional support to improve the attitude estimation.

Finally, temporal and synchronization aspects also play a **crucial role**. Synchronization is *decisive* to avoid disparities and inconsistent triangulations. The frame rate and the subsequent update time strongly influence the **accuracy** of the algorithm's estimation, since excessively low values reduce its quality.

This approach, entirely based on stereo vision and ArUco markers, enables precise attitude estimation. Its accuracy is determined by the quality of calibrations and the robustness of estimation algorithms, but it reaches its maximum effectiveness when **integrated with IMU data**, compensating for the weaknesses of each source.

1.4 Fiducial markers and ArUco tags

The **ArUco marker** is a square planar pattern with a high-contrast black border and an internal matrix encoding a *unique identifier*. This structure allows for fast localization and recognition of the four corners, providing a stable 2D–3D correspondence for attitude estimation. The use of a single marker reduces environmental and economic requirements. However, relying on just one marker requires greater care in localization, calibration, and validation of the measurements.

The recognition pipeline adopts a pair of calibrated cameras. After calibration of the individual cameras and subsequent stereo calibration (*rotation, translation, and epipolar rectification*), the images are acquired and rectified, thus aligning the epipolar lines. The detection of the marker involves several steps:

1. Adaptive thresholding
2. Contour extraction
3. Quadrilateral selection
4. ID decoding

Knowing the marker size, the attitude is calculated through the **PnP algorithm (Perspective-n-Point)**, obtaining the rotation and translation of the reference camera frame. If both cameras observe the marker, the two

estimates can be fused, or the one with higher quality can be selected.

The use of a system based on **stereo vision** improves disparity estimation in the regions close to the marker and enables *metric triangulation* of depth, enhancing stability compared to a monocular approach. This becomes particularly evident when the marker is very small (i.e., far away) or when it is not perpendicular to the viewing angle, as both conditions reduce perspective information.

For each estimate, quality indicators are evaluated: reprojection error, effective presence of the marker (all four corners detected), and the angle between the camera and the marker, penalizing views that are too oblique. Furthermore, a comparison is made between the apparent size of the marker in pixels, in order to identify discrepancies due to triangulation or temporal misalignment.

From a practical point of view, the use of a single ArUco marker offers a **simple and easily reproducible implementation path**: a single print on a rigid surface with a matte finish (to reduce reflections) is sufficient. Under these conditions, a single ArUco marker used in a well-calibrated stereo vision pipeline provides roll, pitch, and yaw measurements while maintaining a *good trade-off between precision and computational demand*.




Figure 1.2: Image caption

1.4.1 Sub-section title

Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas (see Listing 1.1). Suspendisse arcu magna, faucibus ut tincidunt non, ultrices ut turpis. Nullam tristique vehicula massa, id commodo orci sollicitudin vel. Donec nibh ante, ultrices non facilisis sed, mattis id ligula. Sed sed orci sit amet nulla egestas gravida. Suspendisse laoreet, massa vel sagittis gravida, lectus ligula feugiat risus, a aliquam dolor eros ac orci. Nulla egestas tortor quis nunc scelerisque sed tincidunt massa scelerisque. Pellentesque vulputate pharetra lectus, vitae ultricies nisi luctus eu. Nam congue dui eu quam euismod vitae fermentum sem vehicula. Etiam ac leo id nisi placerat posuere. Curabitur mattis augue eget dolor tempus accumsan consequat diam imperdiet. Sed tristique orci id lacus vulputate rhoncus. Morbi tincidunt ante sed turpis luctus tincidunt et sit amet augue. Cum sociis



natoque penatibus et magnis dis parturient montes, nascetur ridiculus mus. Vestibulum ante ipsum primis in faucibus orci luctus et ultrices posuere cubilia Curae; Nunc viverra urna non libero sodales euismod et eleifend sapien. Donec aliquet risus non massa dignissim sollicitudin. Integer a ligula eros. Morbi et lacinia augue [?].

<p>

Pellentesque ac tortor eget eros iaculis euismod
vitae vitae augue.

</p>

<!-- comment -->

Listing 1.1: caption text

Inserimento bibliografico [2] [1] [3]

Bibliography

- [1] E. W. Weisstein, “Euler angles,” *<https://mathworld.wolfram.com/>*, 2009.
- [2] Q. Zhou, Z. Li, G. Yu, H. Li, and N. Zhang, “A novel adaptive kalman filter for euler-angle-based mems imu/magnetometer attitude estimation,” vol. 32, no. 4, p. 045104.
- [3] M. B. Alatisse and G. P. Hancke, “Pose estimation of a mobile robot based on fusion of imu data and vision data using an extended kalman filter,” *Sensors*, vol. 17, no. 10, p. 2164, 2017.